

Kalshiman

The full breakdown - what it is, what it has actually done, and why the plan might work

Data as of 2026-08-20 01:33 UTC

The honest headline: **not profitable yet**. All-time net is -\$77.85 on 291 settled markets at a 79% win rate - small wins, big losses. The last 100 settlements are +\$3.25, and every recent change is aimed at the exact losses that caused the bleeding. Read page 2 before you judge it, and page 9 for the straight answers.

Real settled markets	291
Win rate (all-time)	79% (229W / 62L)
Net P&L (all-time)	-\$77.85
Last 100 settlements	+\$3.25
Last 50 settlements	+\$3.50
Avg win / loss per contract	+12.3c / -48.5c
Model prediction accuracy	90%

Companion markdown lives in docs/ (start with 00-thesis.md and 06-performance.md). Nothing here is financial advice; Kalshi is not available in all states.

1. Is it profitable?

Short answer: **no, not yet** - and here is exactly why, with the account record as the evidence.

All-time record (291 real settlements):

Metric	Value
Settlements	291
Wins / Losses	229 / 62
Win rate	78.7%
Net P&L	-\$77.85
Max drawdown	-\$91.43
Avg win / loss per contract	+12.3c / -48.5c
Median loss (per settlement)	-100.7c (a full reversal)

The math that matters - break-even win rate:

With +12.3c average wins and -48.5c average losses, break-even requires a 134.2% win rate. The all-time rate is 78.7% - just under break-even. That one point is the entire story of the first weeks: the edge (winning 79% of the time) was real, but the losses were too big for the wins.

The turnaround, in 50-settlement windows:


Window	Net	Wins
Settlements 1-50	-\$43.92	33
Settlements 51-100	-\$35.87	37
Settlements 101-150	-\$3.60	43
Settlements 151-200	+\$3.46	46
Settlements 201-250	-\$0.08	39
Settlements 251-291	+\$2.16	31

The system bled ~\$80 in its first 100 settlements (oversized positions, banned market types, whipsaw stops), then stabilized: the last 100 are +\$3.25 and the last 50 are +\$3.50. Recent is a start, not proof - the fixes are only hours old in sample terms.

2. The arc: cumulative P&L;


This is the whole journey in one line. The account sank to roughly -\$90, then the fixes started flattening the curve. The cumulative line is every real settlement in time order (cents).

Cumulative P&L (cents) - all 291 real settlements



The same story as bars - net P&L; per 50-settlement window:


Net P&L by 50-settlement window (cents)



3. The asymmetry (the thing being fixed)

The system wins often and small, and loses rarely and big. The distribution of all 291 settlements tells it plainly: a giant cluster of tiny wins in the 0-20c band, and a tail of catastrophic losses below -100c.


Per-settlement P&L distribution (count by band)



130 settlements landed in the 0-20c band. 34 landed below -100c. The recent fixes attack the left tail directly: \$1-2 position caps, hard bans on the markets that produced the big losses, hold-to-resolution instead of whipsaw stops, and learned-hold-rate confidence in the final minutes.

Win vs loss per contract:

Average per-contract result (cents)



Break-even at these averages needs a 134.2% win rate; the all-time rate is 78.7%. Narrowing that gap is the definition of success for the next 300 settlements.

4. Where the money went (per series)

The losses are not random - they concentrate in identifiable series. Fixing these five rows is the roadmap.

Series	Trades	W/L	Net	Win/c	Loss/c
KXGOLDH	27	24/3	+\$27.24	+14.9c	-75.5c
KXSILVER15M	45	38/7	+\$26.97	+17.1c	-33.9c
KXXRP15M	3	3/0	+\$3.56	+23.1c	0.0c
KXTECHRANKLISTAICODE	1	1/0	+\$1.01	+2.8c	0.0c
KXSILVERW	1	1/0	+\$0.79	+3.0c	0.0c
KXGOLD15M	110	83/27	-\$80.40	+14.9c	-53.0c
KXAAAGASD	1	0/1	-\$25.00	+0.0c	-15.6c
KXMVECROSSCATEGORY-S2026D7AC8BD27AC	1	0/1	-\$10.00	+0.0c	-33.1c
KXWTIH	1	0/1	-\$10.00	+0.0c	-62.7c
KXETH15M	15	9/6	-\$8.52	+6.6c	-67.9c
KXBTC15M	17	14/3	-\$1.91	+5.2c	-32.5c

KXGOLD15M is the leak: -\$80 across 110 trades at a 75% win rate. Its average loss was -\$5.46 - oversized positions and full reversals. Every fix (hold to resolution, \$1-2 caps, learned-hold confidence) targets gold 15m specifically. The single-trade disasters (gas -\$25, combo -\$10, WTI -\$10) are now impossible: those market types are banned and positions are capped.

Net P&L by series (cents) - winners green, losers red



5. The edge, and how the fixes map to the losses

The system does not predict prices; it harvests convergence. In a 15-minute up/down market, the favored side settles correctly 96-99% of the time:

Series	Held	Count	Hold rate
KXSILVERH	569	569	100.0%
KXGOLDH	596	596	100.0%
KXBTCD	34	34	100.0%
KXZEC15M	3596	3616	99.4%
KXWTI15M	1568	1580	99.2%
KXXRP15M	521	525	99.2%
KXDOGE15M	2877	2900	99.2%
KXETH15M	3459	3494	99.0%
KXSILVER15M	8103	8225	98.5%
KXSOL15M	3580	3645	98.2%
KXGOLD15M	7310	7528	97.1%

If the favored side settles your way ~97% of the time and you pay 90-96c for it, each win pays +4-10c and each loss costs -90-96c. That is +EV only if the loss rate stays near the hold-rate - which is why the entire recent effort is about making the losses smaller and rarer.

Loss -> fix mapping (the roadmap):

Loss	Fix (already shipped)
-25c gas blow-up	position caps: every 15m entry is \$1-2
-10c combo	combos banned
-10c WTI close-at-time	exact-spot-at-a-time markets banned
-\$80 gold 15m	65c catastrophe stop, \$1-2 caps, learned-hold confidence in final minutes
-48.5c avg loss per contract	bank at 99.9c, final-minute snipe raises avg win
only 79% win rate (vs 79.6% break-even)	the fixes above pull wins up and losses down until the margin is positive

6. Architecture and strategy in brief

A headless watcher (scripts/watch-headless.js) runs 24/7 under an overnight watchdog. Every 5 seconds it scans enabled strategies plus all eligible closing markets, runs each through four gates (signal, price, timing, risk) plus a gap-risk entry guard, places marketable IOC orders through Kalshi's signed API, manages positions (bank at 99.9c, hold to resolution), settles on resolution, updates per-series learning, and emits Discord alerts.

Entry tiers:

Tier	Condition	Entry	Result
Early-conviction	model + learned $\geq 95\%$	70-98c	65c stop, \$1
Near-certain	model + market in band	80-98c	65c stop, \$1
Final-minute snipe	$\leq 5m$ left, $\max(\text{model}, \text{learned})$ conf	$\leq 98.5c$	\$1
Sure-thing	book $\geq 96.5\%$ implied, $EV > 0$	$\leq 98.5c$	\$1
Weather	3-source forecast + station	80-97c	stop allowed

Models: logistic regression on 9 candle features (implied, deltas, spread, volume rate, OI delta, range, hour), walk-forward backtested, plus the learned hold-rate per series. Prediction accuracy: 90% (154/168).

Markets: gold/silver 15m and hourly, crypto 15m (BTC, ETH, SOL, DOGE, ZEC, BNB, HYPE, NEAR), weather. Hard-banned: exact-spot-at-a-time markets (crypto range/above-below at a date, index-at-close, WTI close-at-time) and combos. Up to 20 concurrent live positions, 20 same-direction.

7. Risk controls

Control	Value
Position size (15m lane)	\$1 (forced); gold $\geq 93c$ \$2 cap
Concurrent live / same-direction	20 / 20
Entry caps	crypto 90c, gold 95c, other 93c (98.5c confirmed)
Stops (15m)	65c catastrophe-only
Bank level	99.9c
Volume minimum	\$1,000
Daily kill switch	halts live trading on loss limit
Re-entry ban	10 min after a stop
Leverage	none on the event lane

8. The lesson that changed everything: stops

Five consecutive trade reviews said the same thing: the stop fired on a whipsaw and the market recovered to WIN. The stop was the leak, not the entry.

Trade	Stop cost	Would have been
SOL NO 89c	-24c	+11c
Silver NO 77c	-33c	+23c
Gold NO 74c	-31c	win
Silver YES 82c	-46c	+18c
Gold YES 83c	-41c	+17c

The favored side settles 96-99% of the time, so a 20-46c whipsaw is noise. Stops converted ~20-30% of winners into losses to protect against a 2-4% tail. The correct protection is now a 65c catastrophe stop plus the entry cap and \$1-2 size: whipsaws recover, and only a true full reversal trips the -65c cap. The last-50 window is positive for the first sustained stretch.

9. Other hard lessons (with receipts)

1. Exact-spot-price-at-a-time markets are near-random - banned (gas -\$25, WTI -\$10).
2. Combos are all-or-nothing - banned (one -\$10 combo).
3. The learned hold-rate beats the live model in the final minutes.
4. Banking early leaves money on the table - bank at 99.9c.
5. Gap-throughs cannot be stopped - cap entry price, keep size small.
6. Crypto tails are fat - 90c cap, tight timing, ETH/XRP handled carefully.
7. Oversized positions cause blow-ups - \$1-2 caps everywhere.
8. The dashboard must match Kalshi's real settlements (now reconciled).
9. Oil is too volatile - WTI banned.
10. Throughput caps were throttling the funnel - uncapped to 20/20.

10. Straight answers

Q: What is the actual edge?

A: Ultra-short 15m markets converge: a 90%+ implied side settles correctly ~96-99% of the time. V2 enters late with a GBM high-confidence gate, rides winners through noise, and caps the rare wrong call at a fixed stop level.

Q: Why did the old system fail?

A: It was a static function: no path/context awareness, no noise-vs-signal sense, paper stops gaped 17-65c, and it never reviewed its own trades. The rework fixed all of it (docs/13-versions.md).

Q: What is the LLM for?

A: A reasoning observer: it reads each round's snapshot like a trader would and gives a YES/NO/PASS opinion with a reason. Advisory only, logged and scored; it never blocks entries (docs/14-reasoning.md).

Q: What are realistic returns?

A: Paper standard hour: 8W/0L, +88.96c, avg +11.12c/trade, \$0 drawdown. At 8 trades/hr x 16 hrs/day x 10c that is ~\$12.80/day paper (docs/15-returns.md).

Q: What are the biggest risks?

A: A full reversal on a high-price entry (bounded by tiered limit-stops to 15c/8c), regime change drifting hold-rates (monitored by review-on-every-close), Kalshi availability/fee changes, and a single perfect hour not being a track record.


Q: Does it use leverage?

A: No. No martingale, no revenge trading, no extreme leverage, no both-sides arbs. Real money is off until live mode is proven.


Q: Why run it at all?

A: A few watts for a self-correcting research system with a quantified edge, capped downside, and honest evaluation of every decision.


Charts




Rework progression: hourly paper P&L; vs avg/trade




Cumulative paper P&L; across the rework hours




V0 vs V1 vs V2 intelligence comparison



GBM out-of-sample accuracy




Predicted vs actual by decile (test set)



The 15-stop review: recovered vs genuinely wrong



Examples of returns by throughput



Doubling vs linear over 30 days

Disclaimer: Past performance is not a guarantee of future results. Trading involves real risk. Kalshi is not available in all states. Nothing in this document is financial advice. All figures are read live from the system state and Kalshi settlement records on 2026-08-20 01:33 UTC.

Sources: memory/headless-state.json (engine state), memory/analysis-data.json (computed from Kalshi /portfolio/settlements), docs/*.md + docs/charts/*.png.