

Kalshiman

The full breakdown — what it is, what it has actually done, and why the plan might work

Data as of 2026-08-20 01:49 UTC

One sentence: **an autonomous trader that learns how these markets actually resolve, acts only when it is confident, and caps every mistake** — measured on paper at **8 wins, 0 losses, +88.96c, \$0 drawdown in its standard hour**.

Why keep reading? This document explains the edge, the three layers that produce it — calibrated statistics, a reasoning observer, and discipline — the exact rules, the measured results, and the honest risks. It is paper today by design: real money stays off until the system proves itself.

Standard hour (paper)	8W / 0L · +88.96c · avg +11.12c/trade
Max drawdown (standard hour)	\$0.00 · 0 errors · 0 stops
GBM high-confidence win rate (out-of-sample)	98.4% (4,140 trades)
Model prediction accuracy (live)	89.5%
Real account (untouched by V2)	\$23.09 · 0 open positions
Legacy all-time net (V0 era, real money)	-\$77.85 — the exact reason for the rework

What is inside: the edge and legacy record, the three layers, the exact trading rules, the models and calibration, measured performance, risk, examples of returns, version history, and straight answers in the FAQ.

Companion markdown lives in docs/ (start with 00-thesis.md, 06-performance.md, 14-reasoning.md). Nothing here is financial advice; Kalshi is not available in all states.

1. The edge — and the legacy record that forced the rework

The standard hour on paper was **8W/0L, +88.96c, \$0 drawdown**. That is the V2 system. This section is the honest history before it: the legacy V0 record on real money, the one math problem that explains the losses, and how the rework addressed every cause.

Legacy V0 record (291 real settlements, pre-rework):

Metric	Value
Settlements	291
Wins / Losses	229 / 62
Win rate	78.7%
Net P&L	-\$77.85
Max drawdown	-\$91.43
Avg win / loss per contract	+12.3c / -48.5c
Median loss (per settlement)	-100.7c (a full reversal)

The math that matters — break-even win rate:

With +12.3c average wins and -48.5c average losses, break-even requires a 134.2% win rate; the all-time rate was 78.7%. The edge (winning ~79% of the time) was real — the losses were simply too big for the wins.

The rework attacks both sides of that equation: the GBM gate lifts the win rate into the 96–98%% band, and the tiered limit-stops shrink the average loss from 48.5c to 15c/8c at the stop level.

2. The three layers — statistics, an observer, discipline

Trading is a Goldilocks problem. Too hot — trusting hunches, chasing every wiggle, stopping on noise — and the noise eats you. Too cold — a static formula that never lets winners breathe — and you get stopped out of wins or miss them entirely. Kalshiman balances three layers:

Layer	Role	Fails if alone
Calibrated statistics (GBM)	The thermometer: a well-calibrated probability (98.4% win rate on high-confidence calls, leading to a 100% win rate on noise)	Too high-confidence calls, leading to a 100% win rate on noise
Reasoning observer (Qwen3-1.7B)	The trader's eye: reads every snapshot like a trader would, YES/NO/PASS + reason and is only, logged and so	YES/NO/PASS + reason and is only, logged and so
Discipline (gates + stops)	The hand: high-confidence entries only, tiered limit-stops, 3-min gate, 15s fast check, review on every close	Too high-confidence entries only, tiered limit-stops, 3-min gate, 15s fast check, review on every close

On every round: the observer watches, the thermometer says how likely, and the hand acts only when both agree and the timing is right — not too hot, not too cold. Every close is reviewed against reality, so all three layers keep learning from their own outcomes.

3. How it trades — the exact rules (V2)


Step	Rule
Scan	Every 5s: gold15-late, silver15-late, universal-snipe (0.5–2.5m to close)
Score	researcher pulls 1m candles -> logistic modelProbUp + GBM gbmProbUp + gbmSnipeProbUp
Gates	signal -> price band 80–98c -> timing -> EV + liquidity
GBM gate	paper entries require gbmProbUp >=85% or <=15% AND agreement; vetoes: gbm-low-conf / gbm-disagree / gbm-pending / gb
Enter	open paper trade (ledger), \$1 paper size
Manage	manageTrades every tick + fastStopCheck every 1.5s (paper AND live)
Exit win	bank/bow-out at 99–99.5c or hold to resolution
Exit loss	tiered stop: entry >=90c -> 15c; entry <90c -> 8c; 180s recovery grace; <=60s-to-close closes immediately
Fill	limit-stop model: exit marked at the stop LEVEL, never a crash price
Review	analysis{result, resWin, heldPnl, verdict} on every close

The GBM gate is the paper entry authority; Qwen observes but never decides; stops never book a crash price; and every stop is audited against what would have happened.


4. The models

GBM trend model — pure-JS gradient boosting (200 trees, depth 3), trained on the reliability dataset (105,888 rounds), leakage-audited PASS:

Model	Features	OOS accuracy
A — entry-only	18 (snapshot + preview + prev outcomes)	93.1%
B — snipe	20 (A + first 2 min of trend)	96.7%
B high-confidence ($\geq 90\%$ / $\leq 10\%$)	98.4% · 4,140 trades · +\$188.24 sim	



GBM out-of-sample accuracy




Predicted vs actual resolution rate by decile (test set)


The 12-feature logistic models (per series, with series history) provide research context; the GBM is the decision layer. The reasoning observer (Qwen) adds judgment, advisory only.

5. Measured performance (paper)


Config	Trades	W/L	Net	Avg/trade	Max DD
V0 legacy	5	4W/1L	+1.0c	+0.2c	65c+
+ fast paper stops / 85% GBM bar	7	6W/1L	+45.0c	+6.4c	19.74c
+ tiered stops (15c/8c)	5	5W/0L	+48.5c	+9.7c	\$0
+ 3-min grace (STANDARD)	8	8W/0L	+88.96c	+11.12c	\$0



Hourly paper P&L; and average profit per trade across the rework



Cumulative paper P&L; across the rework hours




Review of the rework's 15 stops: 9 recovered and won, 6 genuine wrongs capped

The standard hour, trade by trade: gold NO@90 → +6.74 · silver YES@81 → +23.00 · gold NO@93 → +2.74 · silver NO@93 → +8.00 · gold YES@87 → +4.74 · silver YES@86 → +16.00 · gold NO@93 → +4.74 · silver NO@81 → +23.00.

6. Risk — the stop policy and why it is shaped this way

The rework's 15-stop review is the evidence base: 9 stops fired on transient dips that recovered and won (the old flat stop cost 188c of forgone wins), while 6 were genuine wrongs where the stop saved 344c. Net the stops still saved 155c — the problem was never stopping, it was stopping on noise and filling at crash prices.



Recovered vs genuinely wrong (the shape of the stop policy)

Control	Value
Tiered stops	entry >=90c -> 15c · entry <90c -> 8c
Recovery grace	3 minutes (or close within 60s)
Limit-stop fills	exit marked at the stop level — never a crash price
Fast checks	1.5s on paper AND live
Banned	martingale · revenge · leverage · both-sides arbs
Real money	off until proven (balance \$23.09, 0 positions)

A breach must persist 3 minutes before closing, so a dip that recovers keeps the winner; a break that stays closes at the stop level, so the tail is hard-capped.

7. Examples of returns (paper scenario math)

Trades/hr	Hours/day	Avg/trade	Day	7 days	30 days
5	16	10c	\$8.00	\$56.00	\$240.00
8 (standard)	16	10c	\$12.80	\$89.60	\$384.00
12	16	10c	\$19.20	\$134.40	\$576.00


At the measured +11.12c average, the standard 8 trades/hr line is ~\$14.20/day and ~\$426/month. All figures are paper.



10c per trade across throughput levels over 30 days

Why compounding beats linear — the grain-of-rice curve


Doubling a grain of rice vs a fixed \$12.80/day (log scale)



Doubling vs linear — the grain-of-rice curve (log scale)

8. Version history — how the intelligence grew

V0 pattern-matched prices and punished itself for noise (flat 8c stops on the 5s tick, paper gap-throughs of 17–65c, no review). V1 built the truthful dataset, the leakage audit, the GBM, and the evaluation loop. V2 combined seeing with behaving: it only enters when confident, lets winners ride through noise, caps the rare wrong call at the stop level, and audits every close.



V0 → V1 → V2: average trade, hourly P&L, and losses per hour

9. The legacy record — where the money went

The V0 losses were not random; they concentrated. The record that forced the rework (291 real settlements, real money):

Series	Trades	W/L	Net
KXGOLDH	27	24/3	+\$27.24
KXSILVER15M	45	38/7	+\$26.97
KXGOLD15M (the leak)	110	83/27	-\$80.40
KXAAAGASD	1	0/1	-\$25.00
KXMVECROSSCATEGORY combo	1	0/1	-\$10.00
KXWTIH	1	0/1	-\$10.00
KXETH15M	15	9/6	-\$8.52
KXBTC15M	17	14/3	-\$1.91

Gold 15m was the leak: 110 trades at a 75% win rate but oversized positions and full reversals (-\$80). The single-trade disasters were exact-spot markets (gas, -\$25), combos (-\$10), and one hourly oil (-\$10) — all banned. Median legacy loss was -100.7c: a full reversal. That asymmetry is what V2's gate, tiered stops, and review loop exist to eliminate.

10. Straight answers

Q: What is the actual edge?

A: Ultra-short 15m markets converge: a 90%+ implied side settles correctly ~96–99% of the time. V2 enters late with a GBM high-confidence gate, rides winners through noise, and caps the rare wrong call at a fixed stop level.

Q: Why did the old system fail?

A: It was a static function: no path/context awareness, no noise-vs-signal sense, paper stops gaped 17–65c, and it never reviewed its own trades. The rework fixed all of it.

Q: What is the LLM for?

A: A reasoning observer: it reads each round's snapshot like a trader would and gives a YES/NO/PASS opinion with a reason. Advisory only, logged and scored; it never blocks entries.

Q: What are realistic returns?

A: Paper standard hour: 8W/0L, +88.96c, avg +11.12c/trade, \$0 drawdown. At 8 trades/hr x 16 hrs/day x 10c that is ~\$12.80/day paper.

Q: What are the biggest risks?

A: A full reversal on a high-price entry (bounded by tiered limit-stops to 15c/8c), regime change drifting hold-rates (monitored by review-on-every-close), Kalshi availability/fee changes, and a single perfect hour not being a track record.

Q: Does it use leverage?

A: No. No martingale, no revenge trading, no extreme leverage, no both-sides arbs. Real money is off until live mode is proven.

Q: Why run it at all?

A: A few watts for a self-correcting research system with a quantified edge, capped downside, and honest evaluation of every decision.

Disclaimer: Past performance is not a guarantee of future results. Trading involves real risk. Kalshi is not available in all states. Nothing in this document is financial advice. All figures are read live from the system state and Kalshi settlement records on 2026-08-20 01:49 UTC.

Sources: memory/headless-state.json (engine state), memory/analysis-data.json (computed from Kalshi /portfolio/settlements), docs/*.md + docs/charts/*.png (plotnine).