

Kalshiman

Investment brief - an edge engine for prediction markets, validated by replay
and paper

August 20, 2026 | All figures from the live engine unless noted

Confidential - for discussion with prospective investors

1. The thesis

Kalshiman is a software edge in ultra-short prediction markets. It buys the favorite at 90-98 cents when a 15-minute market is near-certain to resolve in its favor, banks the win at 99.5 cents, and caps the rare wrong call with a hard stop and a final-minute escape hatch. The edge is a percentage of position, not a flat number: an 11% win on a \$1 stake is eleven cents, on a \$50 stake it is five dollars and fifty-six cents. Position size is the multiplier; compounding is the engine.

The unit of edge is a percent, not a cent. The system that pays cents on \$1 pays dollars on \$50 - the percentage never changes, the dollars do.

The engine was rebuilt this week after the first version bled. The rework cut the loss-making lanes, tuned the exit hatch to its measured optimum, and added a model plus learned-history signal gate. The result below is the faithful replay of the new engine against 105,888 real market rounds, plus the same engine running live on paper against real Kalshi markets today.

2. The numbers

Faithful replay of the shipped engine (full dataset, no cherry-picking):

Metric	Value
Trades	608
Win rate	77.3% (470W / 138L)
Net P&L	+\$6.29
Average per trade	+1.03 cents
Max drawdown	\$1.96
Exit mix	452 bow-outs / 117 holds / 37 hatches / 2 stops




Figure 1 - cumulative replay equity, cents

The same engine, live on paper today (real Kalshi markets, \$1-3 paper stakes):

Window	Trades	Win rate	Net	Avg/trade
Since rework deployed (Aug 20)	59	88.1% (52W/7L)	+\$1.57	+2.7c
Since silver lane re-enabled	45	95.6% (43W/2L)	+\$1.97	+4.4c
Gold 15m (day)	37	83.8% (31W/6L)	+\$0.86	+2.3c
Silver 15m (day)	21	95.2% (20W/1L)	+\$0.69	+3.3c

Prediction accuracy: model 90.0% vs market 98.5% (200 calls). Fees: 0.63c per contract entry, measured from real fills. Account state: \$23.09 event lane plus \$4.78 margin lane, zero open risk. No real-money trades with the new engine yet - paper validation is the current phase by design.

3. Where the edge comes from




Figure 2 - replay edge by lane, cents

The rework was a subtraction problem. Hourly gold and silver, plus silver 15m, produced the entire historical loss (-\$7,965 of it on the replay). The kept lanes - gold 15m, gold daily, silver daily - produced the entire profit. The engine now trades the kept lanes and skips the rest by rule, not by hope.

4. Why it compounds




Figure 3 - illustrative compounding: 93% win rate, 25% stake, +11% wins / -8% stops

The per-trade return is a percentage of position, and position is a percentage of balance. At 93% wins with a 25% stake ratio, balance grows about +2.4% per trade - \$200 to \$500 in roughly 39 trades. The same math that pays cents at \$1 pays dollars at \$50, which is why the system does not start scaling stakes until \$200-300 is banked as backing.

Balance	Stake (25%)	One win (+11%)	One stop (-8%)
\$5	\$1.25	+\$0.14	-\$0.11
\$30	\$7.50	+\$0.83	-\$0.67
\$100	\$25	+\$2.78	-\$2.22
\$200	\$50	+\$5.56	-\$4.44
\$500	\$125	+\$13.89	-\$11.11

What caps it - the honest part. Liquidity: a 15-minute gold round carries \$50-100k of depth; a \$500 position is fine, a \$50,000 one moves the book and destroys the edge. Loss tail: at 95.2% win rate, one miss in 21 trades; a full final-minute loss at \$20 stakes costs \$20. Small accounts are fragile below about \$4 because the minimum contract forces near all-in. These are the reasons the projection is a curve, not a promise.

5. The network - how it scales

One position cannot get much larger without hitting the liquidity wall, but the same engine can run across more markets and more machines that share what they learn. Ten workers trading \$20 positions across different rounds compound the way one worker trading \$200 cannot. The pool is the compounding story scaled sideways.

Network size	Trades/day each	Per-bot / month	Network / month
1 bot	100	\$6,095	\$6,095
10 bots	100	\$6,095	\$60,950
100 bots	100	\$6,095	\$609,500
1,000 bots	100	\$6,095	\$6,095,000

Assumes 95.2% win rate, \$20 stakes from \$200 backing, +\$2.03 net per trade with capped stops. The loss tail cuts per-trade net roughly in half, and a thousand bots cannot all trade gold at once - the fleet must spread across hundreds of markets. Adoption and market coverage are the real requirements, which is the point of the network.

6. Risk - stated plainly

The real-money record of the prior system was a loss: 360 settled trades, 72% win rate, net -\$90.19, \$17.49 in fees. That version traded flat tapes, paid taker fees on noise, and had no escape hatch. It is the reason the rework exists and the reason this brief leads with the replay and paper evidence instead of the old record.

Three risks remain: (1) the liquidity wall above certain position sizes; (2) the loss tail - one wrong final-minute hold per ~21 trades; (3) unproven live performance at meaningful size with the new engine. Everything in this brief is simulation and paper until the live record says otherwise. That is deliberate, and it is the most honest sentence in the document.

7. The ask

What capital is for: market coverage (more lanes and markets to trade), the network build-out (the pool that lets workers share signals and scale sideways), and the tooling to run the fleet reliably. What investors get: a share of network economics, milestone-based, with the paper-to-live validation as the first milestone gate.

We need money in to cover more markets and stand up the network. The engine is proven on replay and paper; capital buys the coverage that turns that proof into scale.

Contact: kalshiman project lead (details on request). This brief is a discussion document, not an offer of securities.