

Kalshiman

The machine that harvests near-certain markets - and the plan to deploy it as a fleet

A capital deployment brief for investors who want to own part of a working edge

August 20, 2026 | All figures from the live engine unless noted

1. The machine

Prediction markets are the only venue where a 90%+ favorite appears every 15 minutes, all day, every day. A market asks: will gold be up in the next 15 minutes? The favorite trades at 90-98 cents because it usually wins. Our engine buys the favorite only when the setup is near-certain, banks the win at 99.5 cents, and caps the rare wrong call with a hard stop and a measured escape hatch. The edge is structural: the market structure creates it, the engine captures it, and the rails cap the tail.

We do not predict markets. We harvest the certainty they already price. The engine refuses every setup that is not near-certain - that refusal is the edge.

This is not a thesis that needs the market to go up. It needs the market to keep asking yes/no questions every 15 minutes, which it does, thousands of times a day, on a regulated exchange. The engine was rebuilt this week after the first version bled; the rebuild cut the losing lanes by rule, tuned the exit to its measured optimum, and the result is a replay and a paper record that are both positive for the first time.

2. Proof, at small size

Faithful replay of the shipped engine against 105,888 real market rounds - no cherry-picking:

Metric	Value
Trades	608
Win rate	77.3% (470W / 138L)
Net P&L	+\$6.29
Average per trade	+1.03 cents
Max drawdown	\$1.96
Exit mix	452 bow-outs / 117 holds / 37 hatches / 2 stops




Figure 1 - cumulative replay equity, cents. The same engine, live on paper today: 59 trades, 88% win rate, +\$1.57 on \$1-3 paper stakes - and 95.6% since the silver lane returned.

Small size is where the edge is proven, not where it is meant to stay. The margin of the engine is a percentage of position: an 11% win on \$1 is eleven cents, on \$100 it is eleven dollars. Everything below is the business of scaling that percentage without breaking the market it feeds on.

3. Unit economics - one bot is a business

One bot with \$200-5,000 of backing, trading 80+ times a day at a 95%+ win rate, with losses capped by rule. The per-bot model at graduated stakes (80 trades/day, 10.2% net per trade, capped stops):

Stake	Per-trade net	Day (80 trades)	Month (30 days)
\$20	\$2.03	\$162	\$4,872
\$50	\$5.08	\$406	\$12,192
\$100	\$10.16	\$813	\$24,384
\$500	\$50.80	\$4,064	\$121,920




Figure 2 - per-bot monthly net at graduated stakes, modeled unit economics

The cap is never the math - it is the depth of a single 15-minute market. A \$500 position is fine; a \$50,000 one moves the book. That is why the fleet exists: you scale sideways across markets, not one position into the wall. Ten bots at \$50 stakes across different rounds compound like one bot at \$500 - without breaking the market.

4. The deployment - 20 bots, \$5,000 backing each

The plan this capital funds: \$100,000 deployed as 20 independent nodes, each with \$5,000 of backing, spread across markets and lanes. Nodes share what they learn through the pool - a signal proven by one node is verified and adopted by all, so the fleet learns as one organism and loses as diversified units, not as a correlated bet.

Deployment	Per-bot backing	Stake ramp	Fleet month (at \$100 stakes)
20 nodes	\$5,000	\$20 to \$500, milestone-gated	\$487,680

The ramp is the discipline: paper first, then \$20 live, then \$100, then \$500 - each level must prove profitable before the next opens. That is how conviction and risk control share the same sentence. The capital compresses the ladder from years of organic growth to a funded sprint, and the network is the multiplier at every rung.

5. Why now

Kalshi is a regulated US exchange and the venue is investing in precision: it just moved its flagship markets to sub-cent ticks so engines like ours can price and exit exactly. Prediction markets are the fastest-growing category in finance, and ultra-short markets are the segment built for automation. The edge we harvest today is the edge the venue is actively making easier to harvest.

6. What we watch - stated with the same honesty

The first engine lost \$90 trading noise, and that record is in this document because it is the reason the current engine exists. The rebuild is structural, not cosmetic: flat tapes are refused by rule, fees gate every entry, and the tail is capped. Three risks remain and we watch them daily: liquidity above certain stakes, win-rate decay in a regime change, and venue policy. None of them is a secret, and the fleet is the answer to all three - diversification across markets, lanes, and machines.

7. The deployment ask

We are not raising to keep the lights on. We are raising to buy coverage - more markets, more lanes, more nodes - because coverage is the only thing that scales a per-trade edge. Capital deploys in stages, and investors own a share of fleet economics with milestone gates:

Milestone	Gate	What it unlocks
M1 - Live at \$20 stakes	2 profitable weeks per node	The ladder starts climbing
M2 - \$100 stakes	M1 proven at scale	Per-bot pace above \$24k/month
M3 - \$500 stakes + 20 nodes	M2 proven across markets	Full fleet deployment

The edge is proven at \$1. We are going to prove it at \$20, then \$100, then \$500 - and this capital is what compresses that into a funded sprint instead of a years-long grind.
This is a deployment, not a donation.

Contact: kalshiman project lead (details on request). This brief is a discussion document, not an offer of securities.